



Machine Learning/Data Science Interview Cheat sheets

Aqeel Anwar
Version: 0.1.0.3

This document contains cheat sheets on various topics asked during a Machine Learning/Data science interview. This document is constantly updated to include more topics.

[Click here to get the updated version](#)

Table of Contents

Machine Learning Topics	2
1. Bias-Variance Trade-off	2
2. Imbalanced Data in Classification	3
3. Principal Component Analysis	4
4. Bayes' Theorem and Classifier	5
5. Regression Analysis	6
6. Regularization in ML	7
7. Convolutional Neural Network	8
8. Famous CNNs	9
9. Ensemble Methods in Machine Learning	10
10. Autoencoder and Variational Autoencoder	11
Interview Preparation	12
1. Data structures	12
2. Preparing for Coding Interviews	13
3. How to prepare for behavioral interview?	14
4. How to answer a behavioral question?	16

Cheat Sheet – Bias-Variance Tradeoff

What is Bias?

- Error between average model prediction and ground truth
- The bias of the estimated function tells us the capacity of the underlying model to predict the values

$$bias = \mathbb{E}[f'(x)] - f(x)$$

What is Variance?

- Average variability in the model prediction for the given dataset
- The variance of the estimated function tells you how much the function can adjust to the change in the dataset


$$variance = \mathbb{E}[(f'(x) - \mathbb{E}[f'(x)])^2]$$

High Bias

- Overly-simplified Model
- Under-fitting
- High error on both test and train data

High Variance

- Overly-complex Model
- Over-fitting
- Low error on train data and high on test
- Starts modelling the noise in the input



Bias variance Trade-off

- Increasing bias (not always) reduces variance and vice-versa
- Error = bias² + variance + irreducible error
- The best model is where the error is reduced.
- Compromise between bias and variance

Cheat Sheet – Imbalanced Data in Classification



Accuracy doesn't always give the correct insight about your trained model

Accuracy: %age correct prediction

Precision: Exactness of model

Recall: Completeness of model

F1 Score: Combines Precision/Recall

Correct prediction over total predictions
From the detected cats, how many were actually cats

Correctly detected cats over total cats
Harmonic mean of Precision and Recall

One value for entire network
Each class/label has a value

Each class/label has a value
Each class/label has a value

Performance metrics associated with Class 1

		Actual Labels	
		1	0
Predicted Labels	1	True Positive	False Positive
	0	False Negative	True Negative

(Is your prediction correct?) (What did you predict)

True Negative

(Your prediction is correct) (You predicted 0)

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{False +ve rate} = \frac{\text{FP}}{\text{TN} + \text{FP}}$$

$$\text{F1 score} = 2 \times \frac{(\text{Prec} \times \text{Rec})}{(\text{Prec} + \text{Rec})}$$

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{FN} + \text{FP} + \text{TN}}$$


$$\text{Specificity} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

$$\text{Recall, Sensitivity} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$


True +ve rate

Possible solutions


- Data Replication:** Replicate the available data until the number of samples are comparable
- Synthetic Data:** Images: Rotate, dilate, crop, add noise to existing input images and create new data
- Modified Loss:** Modify the loss to reflect greater error when misclassifying smaller sample set
- Change the algorithm:** Increase the model/algorithm complexity so that the two classes are perfectly separable (Con: Overfitting)



$$\text{loss} = a * \text{loss}_{\text{green}} + b * \text{loss}_{\text{blue}} \quad a > b$$



Increase model complexity



Cheat Sheet – PCA Dimensionality Reduction

What is PCA?

- Based on the dataset find a new set of orthogonal feature vectors in such a way that the data spread is maximum in the direction of the feature vector (or dimension)
- Rates the feature vector in the decreasing order of data spread (or variance)
- The datapoints have maximum variance in the first feature vector, and minimum variance in the last feature vector
- The variance of the datapoints in the direction of feature vector can be termed as a measure of information in that direction.

Steps

1. Standardize the datapoints
2. Find the covariance matrix from the given datapoints
3. Carry out eigen-value decomposition of the covariance matrix
4. Sort the eigenvalues and eigenvectors

$$X_{new} = \frac{X - \text{mean}(X)}{\text{std}(X)}$$

$$C[i, j] = \text{cov}(x_i, x_j)$$

$$C = V \Sigma V^{-1}$$

$$\Sigma_{\text{sort}} = \text{sort}(\Sigma) \quad V_{\text{sort}} = \text{sort}(V, \Sigma_{\text{sort}})$$

Dimensionality Reduction with PCA

- Keep the first m out of n feature vectors rated by PCA. These m vectors will be the best m vectors preserving the maximum information that could have been preserved with m vectors on the given dataset

Steps:

1. Carry out steps 1-4 from above
2. Keep first m feature vectors from the sorted eigenvector matrix
3. Transform the data for the new basis (feature vectors)
4. The importance of the feature vector is proportional to the magnitude of the eigen value

$$V_{\text{reduced}} = V[:, 0 : m]$$

$$X_{\text{reduced}} = X_{\text{new}} \times V_{\text{reduced}}$$

Figure 1




Figure 2




Figure 3




Figure 1: Datapoints with feature vectors as x and y-axis

Figure 2: The cartesian coordinate system is rotated to maximize the standard deviation along any one axis (new feature # 2)

Figure 3: Remove the feature vector with minimum standard deviation of datapoints (new feature # 1) and project the data on new feature # 2

Cheat Sheet – Bayes Theorem and Classifier

What is Bayes' Theorem?

- Describes the probability of an event, based on prior knowledge of conditions that might be related to the event.

$$P(A|B) = \frac{P(B|A)(\text{likelihood}) \times P(A)(\text{prior})}{P(B)(\text{evidence})}$$

- How the probability of an event changes when we have knowledge of another event


$$P(A) \longrightarrow P(A|B)$$

Usually, a better estimate than $P(A)$

Example

- Probability of fire $P(F) = 1\%$
- Probability of smoke $P(S) = 10\%$
- Prob of smoke given there is a fire $P(S|F) = 90\%$
- What is the probability that there is a fire given we see a smoke $P(F|S)$?

$$P(F|S) = \frac{P(S|F) \times P(F)}{P(S)} = \frac{0.9 \times 0.01}{0.1} = 9\%$$



Maximum A Posteriori Probability (MAP) Estimation

The MAP estimate of the random variable y , given that we have observed iid $(x_1, x_2, x_3, \dots)$, is given by. We try to accommodate our prior knowledge when estimating.

$$\hat{y}_{MAP} = \underset{y}{\operatorname{argmax}} P(y) \prod_i P(x_i|y)$$

y that maximizes the product of **prior** and **likelihood**

Maximum Likelihood Estimation (MLE)

The MAP estimate of the random variable y , given that we have observed iid $(x_1, x_2, x_3, \dots)$, is given by. We assume we don't have any prior knowledge of the quantity being estimated.

$$\hat{y}_{MLE} = \underset{y}{\operatorname{argmax}} \prod_i P(x_i|y)$$

y that maximizes only the **likelihood**

MLE is a special case of MAP where our prior is uniform (all values are equally likely)

Naïve Bayes' Classifier (Instantiation of MAP as classifier)

Suppose we have two classes, $y=y_1$ and $y=y_2$. Say we have more than one evidence/features $(x_1, x_2, x_3, \dots)$, using Bayes' theorem

$$P(y|x_1, x_2, x_3, \dots) = \frac{P(x_1, x_2, x_3, \dots | y) \times P(y)}{P(x_1, x_2, x_3, \dots)}$$

Naïve Bayes' theorem assumes the features $(x_1, x_2, \dots)$ are i.i.d. i.e $P(x_1, x_2, x_3, \dots | y) = \prod_i P(x_i | y)$



$$P(y|x_1, x_2, x_3, \dots) = \prod_i P(x_i | y) \frac{P(y)}{P(x_1, x_2, x_3, \dots)}$$

$$\hat{y} = y_1 \text{ if } \frac{P(y_1|x_1, x_2, x_3, \dots)}{P(y_2|x_1, x_2, x_3, \dots)} > 1 \text{ else } \hat{y} = y_2$$

Cheat Sheet – Regression Analysis

What is Regression Analysis?

Fitting a function $f(\cdot)$ to datapoints $y_i=f(x_i)$ under some error function. Based on the estimated function and error, we have the following types of regression



	What does it fit?	Estimated function	Error Function
Linear	A line in n dimensions	$f_{\beta}^{linear}(x_i) = \beta_0 + \beta_1 x_i$	$\sum_{i=0}^m \ y_i - f_{\beta}(x_i)\ ^2$
Polynomial	A polynomial of order k	$f_{\beta}^{poly}(x_i) = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \dots$	$\sum_{i=0}^m \ y_i - f_{\beta}(x_i)\ ^2$
Bayesian Linear	Gaussian distribution for each point		
Ridge	Linear/polynomial		$\sum_{i=0}^m \ y_i - f_{\beta}(x_i)\ ^2 + \sum_{j=0}^n \beta_j^2$
LASSO	Linear/polynomial		$\sum_{i=0}^m \ y_i - f_{\beta}(x_i)\ ^2 + \sum_{j=0}^n \beta_j $
Logistic	Linear/polynomial with sigmoid		

Cheat Sheet – Regularization in ML

What is Regularization in ML?

- Regularization is an approach to address over-fitting in ML.
- Overfitted model fails to generalize estimations on test data
- When the underlying model to be learned is low bias/high variance, or when we have small amount of data, the estimated model is prone to over-fitting.
- Regularization reduces the variance of the model




Figure 1. Overfitting

Types of Regularization:

1. Modify the loss function:

- **L2 Regularization:** Prevents the weights from getting too large (defined by L2 norm). Larger the weights, more complex the model is, more chances of overfitting.

$$loss = error(y, \hat{y}) + \lambda \sum_j \beta_j^2 \quad \lambda \geq 0, \lambda \propto \text{model bias}, \lambda \propto \frac{1}{\text{model variance}}$$

- **L1 Regularization:** Prevents the weights from getting too large (defined by L1 norm). Larger the weights, more complex the model is, more chances of overfitting. L1 regularization introduces sparsity in the weights. It forces more weights to be zero, than reducing the the average magnitude of all weights

$$loss = error(y, \hat{y}) + \lambda \sum_j |\beta_j| \quad \lambda \geq 0, \lambda \propto \text{model bias}, \lambda \propto \frac{1}{\text{model variance}}$$

- **Entropy:** Used for the models that output probability. Forces the probability distribution towards uniform distribution.

$$loss = error(y, \hat{y}) + \lambda \sum_i p_i \ln(p_i) \quad \lambda \geq 0, \lambda \propto \text{model bias}, \lambda \propto \frac{1}{\text{model variance}}$$

2. Modify data sampling:

- **Data augmentation:** Create more data from available data by randomly cropping, dilating, rotating, adding small amount of noise etc.
- **K-fold Cross-validation:** Divide the data into k groups. Train on (k-1) groups and test on 1 group. Try all k possible combinations.

3. Change training approach:

- **Injecting noise:** Add random noise to the weights when they are being learned. It pushes the model to be relatively insensitive to small variations in the weights, hence regularization
- **Dropout:** Generally used for neural networks. Connections between consecutive layers are randomly dropped based on a dropout-ratio and the remaining network is trained in the current iteration. In the next iteration, another set of random connections are dropped.




Figure 2. K-fold CV




Figure 3. Drop-out

Cheat Sheet – Convolutional Neural Network

Convolutional Neural Network:

The data gets into the CNN through the input layer and passes through various hidden layers before getting to the output layer. The output of the network is compared to the actual labels in terms of loss or error. The partial derivatives of this loss w.r.t the trainable weights are calculated, and the weights are updated through one of the various methods using backpropagation.

CNN Template:


Most of the commonly used hidden layers (not all) follow a pattern

1. Layer function: Basic transforming function such as convolutional or fully connected layer.


a. Fully Connected: Linear functions between the input and the

a. Convolutional Layers: These layers are applied to 2D (3D) input feature maps. The trainable weights are a 2D (3D) kernel/filter that moves across the input feature map, generating dot products with the overlapping region of the input feature map.

b. Transposed Convolutional (DeConvolutional) Layer: Usually used to increase the size of the output feature map (Upsampling) The idea behind the transposed convolutional layer is to undo (not exactly) the convolutional layer



Dataset: (x , y)



2. Pooling: Non-trainable layer to change the size of the feature map


a. Max/Average Pooling: Decrease the spatial size of the input layer based on selecting the maximum/average value in receptive field defined by the kernel

b. UnPooling: A non-trainable layer used to increase the spatial size of the input layer based on placing the input pixel at a certain index in the receptive field of the output defined by the kernel.

3. Normalization: Usually used just before the activation functions to limit the unbounded activation from increasing the output layer values too high

a. Local Response Normalization LRN: A non-trainable layer that square-normalizes the pixel values in a feature map within a local neighborhood.

b. Batch Normalization: A trainable approach to normalizing the data by learning scale and shift variable during training.



3. Activation: Introduce non-linearity so CNN can efficiently map non-linear complex mapping.

a. Non-parametric/Static functions: Linear, ReLU



b. Parametric functions: ELU, tanh, sigmoid, Leaky ReLU

c. Bounded functions: tanh, sigmoid

5. Loss function: Quantifies how far off the CNN prediction is from the actual labels.

a. Regression Loss Functions: MAE, MSE, Huber loss

b. Classification Loss Functions: Cross entropy, Hinge loss



Cheat Sheet – Famous CNNs

AlexNet – 2012

Why: AlexNet was born out of the need to improve the results of the ImageNet challenge.

What: The network consists of 5 Convolutional (CONV) layers and 3 Fully Connected (FC) layers. The activation used is the Rectified Linear Unit (ReLU).

How: Data augmentation is carried out to reduce over-fitting, Uses Local response localization.

AlexNet Network - Structural Details												
Input	Output	Layer	Stride	Pad	Kernel size	in	out	# of Param				
227	227	3	55	55	96	conv1	4	0	11	11	3	96
55	55	96	27	27	96	maxpool1	2	0	3	3	96	96
27	27	96	27	27	256	conv2	1	2	5	5	96	256
27	27	256	13	13	256	maxpool2	2	0	3	3	256	256
13	13	256	13	13	384	conv3	1	1	3	3	256	384
13	13	384	13	13	384	conv4	1	1	3	3	384	384
13	13	384	13	13	256	conv5	1	1	3	3	384	256
13	13	256	7	7	256	maxpool5	2	0	3	3	256	256
7	7	256	7	7	256	fc6	1	1	9216	4096	37752832	
7	7	256	7	7	256	fc7	1	1	4096	4096	16781312	
7	7	256	7	7	256	fc8	1	1	4096	1000	4097000	
Total									62,378,344			

VGGNet – 2014

Why: VGGNet was born out of the need to reduce the # of parameters in the CONV layers and improve on training time

What: There are multiple variants of VGGNet (VGG16, VGG19, etc.)

How: The important point to note here is that all the conv kernels are of size 3x3 and maxpool kernels are of size 2x2 with a stride of two.

VGG16 - Structural Details												
#	Input Image	output	Layer	Stride	Kernel	in	out	Param				
1	224	224	3	224	224	64	conv3-64	1	3	3	64	1,192
2	224	224	64	224	224	64	conv3-64	1	3	3	64	36928
3	112	112	64	112	112	64	maxpool	2	2	2	64	64
4	112	112	128	112	112	128	conv3-128	1	3	3	128	73856
5	112	112	128	112	112	128	conv3-128	1	3	3	128	147584
6	56	56	128	56	56	128	maxpool	2	2	2	128	65664
7	56	56	256	56	56	256	conv3-256	1	3	3	256	295168
8	56	56	256	56	56	256	conv3-256	1	3	3	256	590080
9	28	28	256	28	28	256	maxpool	2	2	2	256	256
10	28	28	512	28	28	512	conv3-512	1	3	3	512	1180160
11	28	28	512	28	28	512	conv3-512	1	3	3	512	2359808
12	14	14	512	14	14	512	maxpool	2	2	2	512	512
13	14	14	512	14	14	512	conv3-512	1	3	3	512	2359808
14	14	14	512	14	14	512	conv3-512	1	3	3	512	2359808
15	14	14	512	7	7	512	maxpool	2	2	2	512	512
16	1	1	25088	1	1	4096	fc	1	1	25088	4096	102764544
17	1	1	4096	1	1	4096	fc	1	1	4096	4096	16781312
18	1	1	4096	1	1	1000	fc	1	1	4096	1000	4097000
Total									138,423,208			

ResNet – 2015

Why: Neural Networks are notorious for not being able to find a simpler mapping when it exists. ResNet solves that.

What: There are multiple versions of ResNetXX architectures where 'XX' denotes the number of layers. The most used ones are ResNet50 and ResNet101. Since the vanishing gradient problem was taken care of (more about it in the How part), CNN started to get deeper and deeper

How: ResNet architecture makes use of shortcut connections to solve the vanishing gradient problem. The basic building block of ResNet is a Residual block that is repeated throughout the network.




Figure 1 ResNet Block




Figure 2 Inception Block

Inception – 2014


Why: Larger kernels are preferred for more global features, on the other hand, smaller kernels provide good results in detecting area-specific features. For effective recognition of such a variable-sized feature, we need kernels of different sizes. That is what Inception does.

What: The Inception network architecture consists of several inception modules of the following structure. Each inception module consists of four operations in parallel, 1x1 conv layer, 3x3 conv layer, 5x5 conv layer, max pooling

How: Inception increases the network space from which the best network is to be chosen via training. Each inception module can capture salient features at different levels.

GoogleNet - Structural Details												
	Input Image	output	Layer	Input Stride	Pad	Kernel	in	out	Param			
descriptors (3a)	112 112 64	56 56 64	conv1	2	0.5	3 3 3	64	64	0	9,437		
	112 112 64	56 56 64	conv1	2	0.5	3 3 3	64	64	0	9,437		
	56 56 64	56 56 64	conv1	2	0.5	3 3 3	64	64	0	9,437		
	56 56 64	56 56 64	conv1	2	0.5	3 3 3	64	64	0	9,437		
	56 56 192	28 28 192	maxpool2	2	0.5	3 3 3	192	192	0	0		
	28 28 192	28 28 96	conv2a	1	0	1 1 1	192	96	115,728			
	28 28 96	28 28 16	conv2b	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2c	1	0	1 3 3	192	96	115,728			
	28 28 192	28 28 64	conv2d	1	0	1 1 1	64	64	115,728			
	28 28 96	28 28 32	conv2e	1	0	1 1 1	64	32	2,832			
descriptors (3b)	28 28 192	28 28 96	conv2f	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2g	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2h	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2i	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2j	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2k	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2l	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2m	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2n	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2o	1	0	1 1 1	192	16	3,088			
descriptors (4a)	28 28 192	28 28 96	conv2p	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2q	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2r	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2s	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2t	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2u	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2v	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2w	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2x	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2y	1	0	1 1 1	192	16	3,088			
descriptors (4b)	28 28 192	28 28 96	conv2z	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2aa	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2ab	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ac	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2ad	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ae	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2af	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ag	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2ah	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ai	1	0	1 1 1	192	16	3,088			
descriptors (5a)	28 28 192	28 28 96	conv2aj	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ak	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2al	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2am	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2an	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ao	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2ap	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2aq	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2ar	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2as	1	0	1 1 1	192	16	3,088			
descriptors (5b)	28 28 192	28 28 96	conv2at	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2au	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2av	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2aw	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2ax	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ay	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2az	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ba	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bb	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bc	1	0	1 1 1	192	16	3,088			
descriptors (5c)	28 28 192	28 28 96	conv2bd	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2be	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bf	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bg	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bh	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bi	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bj	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bk	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bl	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bm	1	0	1 1 1	192	16	3,088			
descriptors (5d)	28 28 192	28 28 96	conv2bn	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bo	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bp	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bq	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2br	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bs	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bt	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bu	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bv	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2bw	1	0	1 1 1	192	16	3,088			
descriptors (5e)	28 28 192	28 28 96	conv2bx	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2by	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2bz	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ca	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cb	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2cc	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cd	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ce	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cf	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2cg	1	0	1 1 1	192	16	3,088			
descriptors (5f)	28 28 192	28 28 96	conv2ch	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ci	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cj	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ck	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cl	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2cm	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cn	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2co	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cp	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2cq	1	0	1 1 1	192	16	3,088			
descriptors (5g)	28 28 192	28 28 96	conv2cr	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2cs	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2ct	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2cu	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cv	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2cw	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cx	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2cy	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2cz	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2da	1	0	1 1 1	192	16	3,088			
descriptors (5h)	28 28 192	28 28 96	conv2db	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2dc	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2dd	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2de	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2df	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2dg	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2dh	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2di	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2dj	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2dk	1	0	1 1 1	192	16	3,088			
descriptors (5i)	28 28 192	28 28 96	conv2dl	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2dm	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2dn	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2do	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2dp	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2dq	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2dr	1	0	1 3 3	192	96	115,728			
	28 28 96	28 28 16	conv2ds	1	0	1 1 1	192	16	3,088			
	28 28 192	28 28 96	conv2dt	1	0	1 3 3	192	96	115,728			
	28 28 96	28										

Cheat Sheet – Ensemble Learning in ML





Cheat Sheet – Autoencoder & Variational Autoencoder

Context – Data Compression

- Data compression is an essential phase in training a network. The idea is to compress the data so that the same amount of information can be represented by fewer bits.



Auto Encoder (AE)

- Autoencoder is used to learn efficient embeddings of unlabeled data for a given network configuration. It consists of two parts, an **encoder**, and a **decoder**.
- The encoder compresses the data from a higher-dimensional space to a lower-dimensional space (also called the latent space), while the decoder converts the latent space back to higher-dimensional space.
- The entire encoder-decoder architecture is collectively trained on the loss function which encourages that the input is reconstructed at the output. Hence the loss function is the mean squared error between the encoder input and the decoder output.
- The latent variable is not regularized. Picking a random latent variable will generate garbage output.
- Latent variable is deterministic values and the space lacks the generative capability



Variational Auto Encoder (VAE)

- Variational autoencoder addresses the issue of non-regularized latent space in autoencoder and provides the generative capability to the entire space.
- Instead of outputting the vectors in the latent space, the encoder of VAE outputs parameters of a pre-defined distribution in the latent space for every input.
- The VAE then imposes a constraint on this latent distribution forcing it to be a normal distribution.
- The latent variable in the compressed form is mean and variance
- The training loss of VAE is defined as the sum of the reconstruction loss and the similarity loss (the KL divergence between the unit gaussian and decoder output distribution).
- The latent variable is smooth and continuous i.e., random values of latent variable generates meaningful output at the decoder, hence the latent space has generative capabilities.
- The input of the decoder is sampled from a gaussian with mean/variance of the output of encoder.



Cheat Sheet – Data Structures

1. List

- Ordered collection of elements
- The position of each element is defined by the *index*
- The elements can be accessed in *any order*

	len = 4				
value	12	2	0	6	18
index	0	1	2	3	4





- Each linked list element contains both the values and the address (pointer) to the next linked list element.
- Hence the linked list can only be traversed sequentially going through each element at a time

2. Linked List

- Linked List does not have their order defined by their physical placement in the memory.
- Contiguous elements of the linked list are not placed adjacent to each other in the memory.

3. Stack

- Stack is a sequential data structure which maintains the order of elements as they were inserted in.
- Last In First Out (LIFO) order, which means that the elements can only be accessed in the reverse order as they were inserted into the stack.
- The element to be inserted last, will be the first one to get removed from the stack.
- Push() adds an element at the head of the stack, while pop() removes an element from the head of the stack
- A real-life example of a stack is a stack of kitchen plates





4. Queue

- A queue is a sequential data structure that maintains the order of elements as they were inserted in
- First In First Out (FIFO), the element to be inserted first, will be the first one to get removed from the queue
- Whenever an element is added (Enqueue()) it is added to the end of the queue. On the other hand, element removal (Dequeue()) is done from the front of the queue.
- A real-life example is a check-out line at a grocery store

5. HashTable


- Creates paired assignments (key mapped to values) so the pairs can be accessed in constant time
- For each (key, value) pair, the key is passed through a hash function to create a unique physical address for the value to be stored in the memory.
- Hash function can end up generating the same physical address for different keys. This is called a collision.



- Maintains a hierarchical relation between its elements.
- Root Node – The node at the top of the tree
- Parent Node – Any node that has at least one child
- Child Node – The successor of a parent node is known as a child node. A node can be both a parent and a child node. The root is never a child node.
- Leaf Node – The node which does not have any child node.
- Traversing – Passing through the nodes in a certain order, e.g BFS, DFS

4. Graph

- A graph is a pair of sets (V, E), where V is set of all the vertices, E is set of all edges.
- A neighbor of a node is set of all vertices connected with that node through an edge.
- As opposed to trees, a graph can be cyclic, which means starting from a node and following the edges, you can end up on the same node



Cheat Sheet – Preparing for Coding Interviews

Part 1 – How to prepare for coding interviews?*

- **The timeline:**




Fig. 1 – Preparation Timeline for Coding Interviews

- **Review Data structures and Complexities:**

The following 7 data structures are necessary for the interview, and their time/space complexity

- List/Arrays, Linked List, Hash Table/dictionary, Tree, Graph, Heap, Queue
- Click [here](#) for tutorial.

- **Practice coding questions:**

- Multiple online resources such as [LeetCode.com](#), [InterviewBit.com](#), [HackerRank.com](#) etc.
- Pick one online resource and aim for easy and medium coding questions (approx. 100-150).
- Beginners start preparing 2-3 months before the interview, and intermediates about 1 month.

- **Note:**

- From my personal experience, paid subscription of LeetCode.com was worth it.
- Facebook, Uber, Google and Microsoft tagged question of LeetCode covered almost 90% of the questions asked

Part 2 – How to answer a coding question?*

- **Listen to the question**

The interviewer will explain the question with an example. Note down the important points.

- **Talk about your understanding of the question**

Repeat the question and confirm your understanding. Ask clarifying questions such as

1. Input/Output data type limitations
2. Input size/length limitations
3. Special/Corner cases

- **Discuss your approach**

Walk through how would you approach the problem and ask the interviewer if he agrees with it.

Talk about the data structure you prefer and why. Discuss the solution with the bigger picture.

- **Start coding**

Ask the interviewer if you could start coding. Define useful functions and explain as you write.

Think out loud so the interviewer can evaluate your thought process.

- **Discuss the time and space complexity**

Discuss the time and space complexity in terms of Big O for your coded approach.

- **Optimize the approach**

If your approach is not the most optimized one, the interviewer will hint you a few improvements. Pay attention to hints and try to optimize your code.





Fig. 2 – How to answer a coding question?

***Disclaimer :** The recommendations are based on personal experiences of the author. The mentioned approach and resources might work great for some, but not so much for others.

Source: <https://www.cheatsheets.aqeel-anwar.com>





Keywords

List important keywords that will be populated with your personal stories. Most common keywords are given in the table below

Conflict Resolution	Negotiation	Compromise to achieve goal	Creativity	Flexibility	Convincing
Handling Crisis	Challenging Situation	Working with difficult people	Another team priorities not aligned	Adjust to a colleague style	Take Stand
Handling -ve feedback	Coworker view of you	Working with a deadline	Your strength	Your weakness	Influence Others
Handling failure	Handling unexpected situation	Converting challenge to opportunity	Decision without enough data	Conflict Resolution	Mentorship/Leadership


Stories

- List all the organizations you have been a part of. For example
 - Academia: BSc, MSc, PhD
 - Industry: Jobs, Internship
 - Societies: Cultural, Technical, Sports
- Think of stories from step 1 that can fall into one of the keywords categories. The more stories the better. You should have at least 10-15 stories.
- Create a summary table by assigning multiple keywords to each stories. This will help you filter out the stories when the question asked in the interview. An example can be seen below

Story 1:	[Convincing] [Take Stand] [influence other]
Story 2:	[Mentorship] [Leadership]
Story 3:	[Conflict resolution] [Negotiation]
Story 4:	[decision-without-enough-data]

STAR Format

Write down the stories in the STAR format as explained in the 2/4 part of this cheat sheet. This will help you practice the organization of story in a meaningful way.



Example: “Tell us about a time when you had to convince senior executives”

Situation

Explain the situation and provide necessary context for your story.

S

“I worked as an intern in XYZ company in the summer of 2019. The project details provided to me was elaborative. After some initial brainstorming, and research I realized that the project approach can be modified to make it more efficient in terms of the underlying KPIs. I decided to talk to my manager about it.”

Task

Explain the task and your responsibility in the situation

T

“I had an hour-long call with my manager and explained him in detail the proposed approach and how it could improve the KPIs. I was able to convince him. He asked me if I will be able to present my proposed approach for approval in front of the higher executives. I agreed to it. I was working out of the ABC(city) office and the executives need to fly in from XYZ(city) office.”

Action

Walk through the steps and actions you took to address the issue

A


“I did a quick background check on the executives to know better about their area of expertise so that I can convince them accordingly. I prepared an elaborative 15 slide presentation starting with explaining their approach, moving onto my proposed approach and finally comparing them on preliminary results.

Result

State the outcome of the result of your actions

R

“After some active discussion we were able to establish that the proposed approach was better than the initial one. The executives proposed a few small changes to my approach and really appreciated my stand. At the end of my internship, I was selected among the 3 out of 68 interns who got to meet the senior vice president of the company over lunch.”



Understand, Extract, Map, Select and Apply

Example: “Tell us about a time when you had to convince senior executives”

Understand

Understand the question

Example: A story where I was able to convince my seniors. Maybe they had something in mind, and I had a better approach and tried to convince them

Extract

Extract keywords and tags

Extract useful keywords that encapsulates the gist of the question

Example:

[Convincing], [Creative], [Leadership]

Map

Map the keyword to your stories

Shortlist all the stories that fall under the keywords extracted from previous step

Example:

Story1, Story2, Story3, Story4, ... , Story N

Select

Select the best story

From the shortlisted stories, pick the one that best describes the question and has not been used so far in the interview


Example: Story3

Apply

Apply the STAR method


Apply the STAR method on the selected story to answer the question

Example: See Cheat Sheet 2/3 for details




Summarizing the behavioral interview

How to prepare for the interview



How to answer a question during interview



Follow the Author:

Follow the author for more machine learning/data science content at

- □ Medium: <https://aqeel-anwar.medium.com>
- □ LinkedIn: <https://www.linkedin.com/in/aqeelanwarmalik/>

Feedback:

If you find any error in the cheat sheets, please provide your feedback [here](#)

Version History

- **Version 0.1.0.3 - Dec 25, 2021**
 - Added cheat sheets: *Autoencoder and Variational Autoencoder*
- **Version 0.1.0.2 - May 19, 2021**
 - Added cheat sheets: *Data structures* and *Preparing for Coding Interview*
 - Added tutorial links at the end of each cheat sheet
- **Version 0.1.0.1 - Apr 05, 2021**

Fixed minor typo issues in Baye's Theorem, Regression analysis and Classifier and PCA dimensionality reduction cheat sheets.
- **Version 0.1.0.0 - Mar 30, 2021**

Initial draft with nine basics of ML and two behavioral interview cheat sheets.